Abhinav Cloud

Principal AI & LLM Systems Engineer

Email: **abhinavclouds@gmail.com** | Phone: +1 (555) 481-9920 | Location: Seattle, WA (Remote)

PROFESSIONAL SUMMARY

Principal AI Platform Engineer specializing in high-throughput LLM serving infrastructure, specialized multi-agent orchestrators, and enterprise vector search. Extensive experience optimizing inference latency with vLLM, FlashAttention-2, and TensorRT-LLM.

TECHNICAL PROFICIENCIES

Key Skills: Python • PyTorch • vLLM • FlashAttention • LangChain • Gemini API • FastAPI • Vector Databases • Qdrant • Ray • Triton Inference Server • CUDA

PROFESSIONAL EXPERIENCE

Principal AI Infrastructure Engineer | *Synthetix AI Labs* (2023 - Present)

- Spearheaded enterprise LLM serving cluster serving 80M daily tokens with vLLM and continuous batching on NVIDIA H100s, slashing p99 Time-To-First-Token by 44%.
- Architected specialized multi-agent evaluation orchestrator with bounded retries, structured JSON outputs, and adversarial critic validation.
- Designed low-latency hybrid semantic search pipeline using Qdrant vector indexing and cross-encoder re-ranking handling 3,500 queries/sec.
- Built prompt-injection defense shield sanitizing untrusted candidate documents and neutralizing delimiter hijack payloads.

Senior Machine Learning Systems Engineer | *DeepVision Intelligence* (2019 - 2023)

- Trained and fine-tuned domain-specific transformer models using PyTorch Distributed Data Parallel (DDP).
- Built end-to-end feature store and model evaluation tracking pipeline serving real-time scoring to production microservices.
- Optimized inference throughput by 3.2x via quantization (AWQ / FP8) and custom Triton kernels.

EDUCATION

M.S. in Artificial Intelligence, Stanford University (2019)