Abhinav Cloud

Principal ML Systems & LLM Platform Engineer

Email: **abhinavclouds@gmail.com** | Phone: +1 (555) 481-9920 | Location: Seattle, WA (Remote)

PROFESSIONAL SUMMARY

Machine Learning Systems Engineer specializing in high-throughput LLM serving infrastructure, multi-agent evaluation topologies, and GPU acceleration. Extensive experience deploying vLLM, FlashAttention, and Triton Inference Server at scale.

TECHNICAL PROFICIENCIES

Key Skills: PyTorch ? vLLM ? FlashAttention ? Triton Server ? Python ? CUDA ? LangChain / LlamaIndex ? Vector DBs (Qdrant, Pinecone) ? Docker ? Kubernetes ? Ray

PROFESSIONAL EXPERIENCE

Principal AI Platform Engineer | *Synthetix AI Labs* (2023 - Present)

? Spearheaded enterprise LLM cluster serving 80M daily tokens with vLLM and continuous batching on NVIDIA H100s, slashing p99 TTFT by 44%.

? Architected specialized multi-agent evaluation orchestrator with bounded retries, structured JSON outputs, and adversarial critic validation.

? Designed low-latency hybrid semantic search pipeline using Qdrant vector indexing and cross-encoder re-ranking handling 3,500 queries/sec.

? Built prompt-injection defense shield sanitizing untrusted candidate documents and neutralizing delimiter hijack payloads.

Senior ML Infrastructure Engineer | *DeepVision Intelligence* (2019 - 2023)

? Trained and fine-tuned domain-specific transformer models using PyTorch Distributed Data Parallel (DDP).

? Optimized inference throughput by 3.2x via quantization (AWQ / FP8) and custom Triton kernels.

EDUCATION

M.S. in Artificial Intelligence, Stanford University (2019)